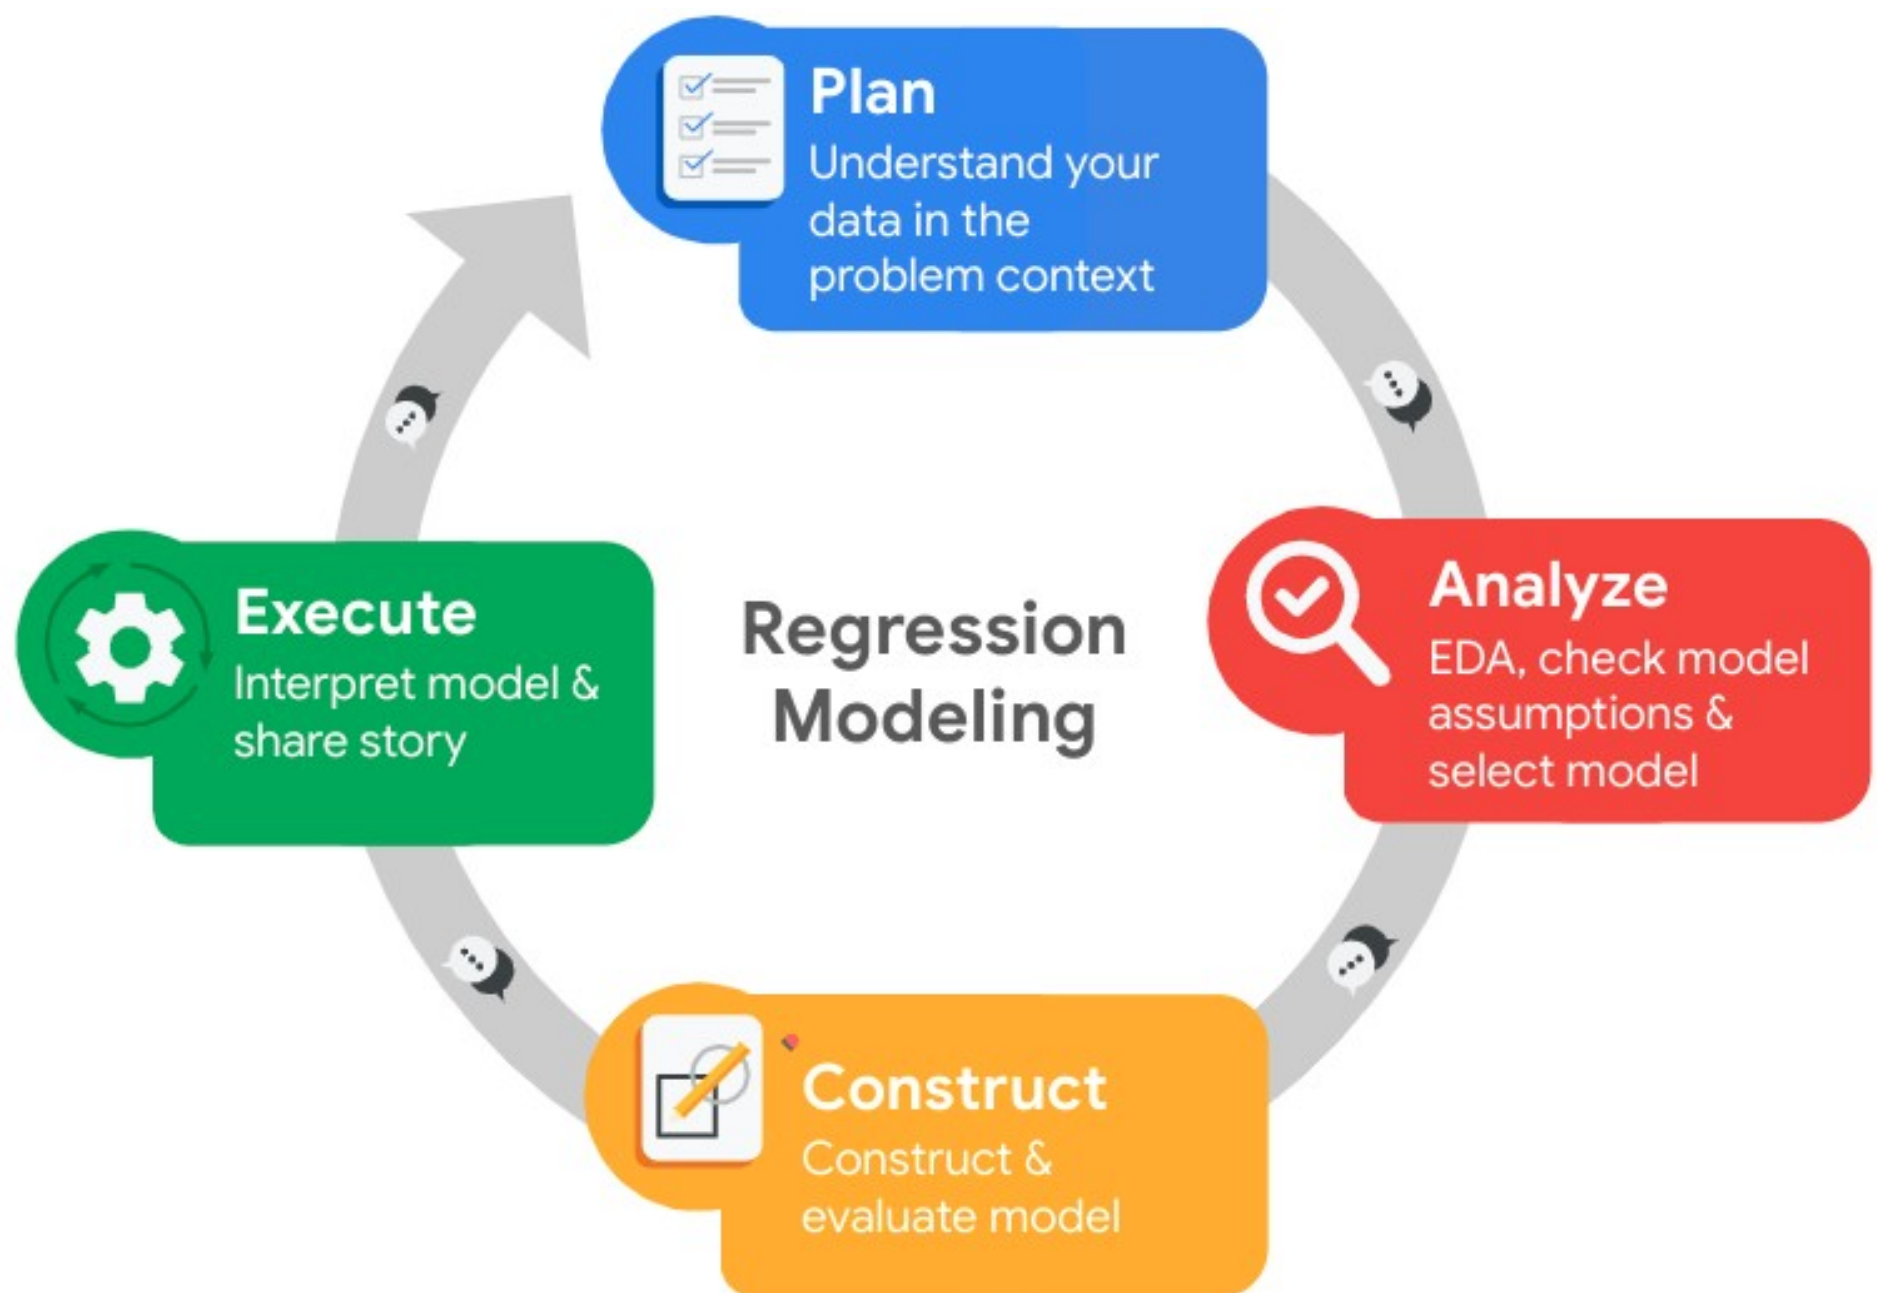


PACE Stages: The Project Framework



Pace: Plan Stage

Case Study Introduction:

The modeling objective is to build and test a classification model that uses banking data to predict whether a customer will churn.

If a customer churns, it means they left the bank and took their business elsewhere. If we can predict customers who are likely to churn, we can take measures to retain them before they do.

These measures could be promotions, discounts, or other incentives to boost customer satisfaction and, therefore, retention.

Our data dictionary shows that there is a column called Exited. This is a Boolean value that indicates whether or not a customer left the bank (0 = did not leave, 1 = did leave)

This will be our target variable. In other words, for each customer, our model should predict whether he should have a 0 or a 1 in the Exited column.

Upon that, We conclude the key considerations in the following points:

Modeling objective: To predict whether a customer will churn—a binary classification task.

Target variable: Exited column—0 or 1.

Class balance: The data is imbalanced 80/20 (not churned/churned).

Primary evaluation metric: F1 score.

Modeling workflow and model selection: The champion model will be the model with the best validation F1 score.

Step 1. Imports

```
import pandas as pd
import numpy as np
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
from imblearn.over_sampling import SMOTE
from collections import Counter
```

```
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import recall_score, precision_score, f1_score, accuracy_score, roc_auc_score,\
    ConfusionMatrixDisplay, confusion_matrix, RocCurveDisplay
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier, plot_importance

import warnings
warnings.filterwarnings('ignore')

# This module lets us save our models once we fit them.
import pickle

# This lets us see all of the columns, preventing Jupyter from redacting them.
pd.set_option('display.max_columns', None)
```

Step 2. Data Exploration (Initial EDA and data cleaning)

Gather basic information about the data

```
original_df = pd.read_csv(r"D:\Google Advanced Data Analytics\Nuts & Bolts of Machine Learning\Churn_Modelling.csv")
original_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber              10000 non-null  int64
1   CustomerId             10000 non-null  int64
2   Surname                10000 non-null  object
3   CreditScore            10000 non-null  int64
4   Geography              10000 non-null  object
5   Gender                 10000 non-null  object
6   Age                   10000 non-null  int64
7   Tenure                 10000 non-null  int64
8   Balance                10000 non-null  float64
9   NumOfProducts          10000 non-null  int64
10  HasCrCard              10000 non-null  int64
11  IsActiveMember         10000 non-null  int64
12  EstimatedSalary        10000 non-null  float64
13  Exited                 10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

Gather descriptive statistics about the data

```
original_df.describe(include= "all")
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	\
count	10000.00000	1.000000e+04	10000	10000.000000	10000	10000	
unique	NaN	NaN	2932	NaN	3	2	
top	NaN	NaN	Smith	NaN	France	Male	
freq	NaN	NaN	32	NaN	5014	5457	
mean	5000.50000	1.569094e+07	NaN	650.528800	NaN	NaN	
std	2886.89568	7.193619e+04	NaN	96.653299	NaN	NaN	
min	1.00000	1.556570e+07	NaN	350.000000	NaN	NaN	
25%	2500.75000	1.562853e+07	NaN	584.000000	NaN	NaN	
50%	5000.50000	1.569074e+07	NaN	652.000000	NaN	NaN	
75%	7500.25000	1.575323e+07	NaN	718.000000	NaN	NaN	
max	10000.00000	1.581569e+07	NaN	850.000000	NaN	NaN	

	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	
unique	NaN	NaN	NaN	NaN	NaN	
top	NaN	NaN	NaN	NaN	NaN	
freq	NaN	NaN	NaN	NaN	NaN	
mean	38.921800	5.012800	76485.889288	1.530200	0.70550	
std	10.487806	2.892174	62397.405202	0.581654	0.45584	
min	18.000000	0.000000	0.000000	1.000000	0.00000	
25%	32.000000	3.000000	0.000000	1.000000	0.00000	
50%	37.000000	5.000000	97198.540000	1.000000	1.00000	
75%	44.000000	7.000000	127644.240000	2.000000	1.00000	
max	92.000000	10.000000	250898.090000	4.000000	1.00000	

	IsActiveMember	EstimatedSalary	Exited
count	10000.000000	10000.000000	10000.000000
unique	NaN	NaN	NaN
top	NaN	NaN	NaN
freq	NaN	NaN	NaN
mean	0.515100	100090.239881	0.203700
std	0.499797	57510.492818	0.402769
min	0.000000	11.580000	0.000000
25%	0.000000	51002.110000	0.000000
50%	1.000000	100193.915000	0.000000
75%	1.000000	149388.247500	0.000000
max	1.000000	199992.480000	1.000000

Check missing values

```
original_df.isna().sum()
```

```
RowNumber      0
CustomerId     0
Surname        0
CreditScore    0
Geography      0
Gender         0
Age           0
Tenure         0
Balance        0
NumOfProducts  0
HasCrCard      0
IsActiveMember 0
EstimatedSalary 0
Exited         0
dtype: int64
```

Check duplicates

```
original_df.duplicated().sum()
```

```
0
```

Check outliers

```
sns.set_style("whitegrid")
sns.set_palette("Dark2")
```

```
fig, ax = plt.subplots(6, 1, figsize= (12, 24))
```

```
sns.boxplot(x= original_df["CreditScore"], showfliers= True, ax= ax[0])
ax[0].set_title("Distribution of Credit Scores", fontsize= 12, fontweight= "bold")
ax[0].set_xlabel("")
```

```
sns.boxplot(x= original_df["Age"], showfliers= True, ax= ax[1])
ax[1].set_title("Distribution of Ages", fontsize= 12, fontweight= "bold")
ax[1].set_xlabel("")
```

```
sns.boxplot(x= original_df["Tenure"], showfliers= True, ax= ax[2])
ax[2].set_title("Distribution of Customers' Tenure", fontsize= 12, fontweight= "bold")
ax[2].set_xlabel("")
```


```
sns.boxplot(x= original_df["Balance"], showfliers= True, ax= ax[3])
ax[3].set_title("Distribution of Customers' Balance", fontsize= 12, fontweight= "bold")
ax[3].set_xlabel("")
```

```
sns.boxplot(x= original_df["NumOfProducts"], showfliers= True, ax= ax[4])
ax[4].set_title("Distribution of Products' Number", fontsize= 12, fontweight= "bold")
ax[4].set_xlabel("")
```


```
sns.boxplot(x= original_df["EstimatedSalary"], showfliers= True, ax= ax[5])
ax[5].set_title("Distribution of Customers' Estimated Salaries", fontsize= 12, fontweight= "bold")
ax[5].set_xlabel("")
```

```
plt.show()
```


Distribution of Credit Scores




Distribution of Ages



Distribution of Customers' Tenure



Distribution of Customers' Balance



Distribution of Products' Number



```
def num_outliers(column_name:str):
    percentile_25 = original_df[column_name].quantile(0.25)
    percentile_75 = original_df[column_name].quantile(0.75)
    iqr = percentile_75 - percentile_25
    upper_limit = percentile_75 + 1.5 * iqr
    lower_limit = percentile_25 - 1.5 * iqr
    num_outliers = len(original_df[(original_df[column_name] > upper_limit) | (original_df[column_name] <
lower_limit)])

    print("Upper Limit =", upper_limit)
    print("Lower Limit =", lower_limit)
    print("Number of outliers =", num_outliers)
```

```
num_outliers("CreditScore")
```

```
Upper Limit = 919.0
Lower Limit = 383.0
Number of outliers = 15
```

```
num_outliers("Age")
```

```
Upper Limit = 62.0
Lower Limit = 14.0
Number of outliers = 359
```

```
num_outliers("NumOfProducts")
```

```
Upper Limit = 3.5
Lower Limit = -0.5
Number of outliers = 60
```

pAce: Analyze Stage

Step 2. Data Exploration (Continue EDA)

```
original_df.groupby(["Exited"])[["CreditScore", "Age", "Tenure", "Balance", "NumOfProducts", "HasCrCard",
"IsActiveMember", "EstimatedSalary"]].median(numeric_only= True)
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
Exited							
0	653.0	36.0	5.0	92072.68	2.0	1.0	
1	646.0	45.0	5.0	109349.29	1.0	1.0	

	IsActiveMember	EstimatedSalary
Exited		
0	1.0	99645.04
1	0.0	102460.84

From the above grouping, we notice the following:

1. on median the churned customer has less credit score with 7 points than the retained customer, given that the credit score ranges between 350 to 850 points.
2. on median the churned customer ages 9 years older than the retained customer.
3. on median the churned customer has almost 18000 dollars on their balance greater than the retained customer.

Data Visualizations 1: Proportions

```
original_df["Exited"].value_counts(normalize= True) * 100
```

```
0    79.63
1    20.37
Name: Exited, dtype: float64
```

```
plt.pie(original_df["Exited"].value_counts(), labels= ["Stayed: 79.63%", "Churned: 20.37%"], explode= [0, 0.2])
my_circle = plt.Circle( (0,0), 0.5, color= 'white')
p= plt.gcf()
p.gca().add_artist(my_circle)
```

```
plt.title("Proportion of Churned Customers:", fontweight= "bold")
plt.show()
```

Proportion of Churned Customers:




```
original_df["Geography"].value_counts(normalize= True) * 100
```

```
France      50.14
Germany     25.09
Spain       24.77
Name: Geography, dtype: float64
```

```
plt.pie(original_df["Geography"].value_counts(), labels= ["France: 50.14%", "Germany: 25.09%", "Spain: 24.77%"])
my_circle = plt.Circle( (0,0), 0.5, color= 'white')
p= plt.gcf()
p.gca().add_artist(my_circle)
```

```
plt.title("Proportion of Each Country:", fontweight= "bold")
plt.show()
```

Proportion of Each Country:




```
original_df["HasCrCard"].value_counts(normalize= True) * 100
```

```
1      70.55
0      29.45
Name: HasCrCard, dtype: float64
```

```
plt.pie(original_df["HasCrCard"].value_counts(), labels= ["Credit Card Holders: 70.55%", "No Credit Card: 29.45%"])
my_circle = plt.Circle( (0,0), 0.5, color= 'white')
p= plt.gcf()
p.gca().add_artist(my_circle)
```

```
plt.title("Proportion of Credit Card Holders:", fontweight= "bold")
plt.show()
```


Proportion of Credit Card Holders:



```
original_df["IsActiveMember"].value_counts(normalize= True) * 100
```

```
1    51.51  
0    48.49
```

```
Name: IsActiveMember, dtype: float64
```

```
plt.pie(original_df["IsActiveMember"].value_counts(), labels= ["Active Members: 51.51%", "Idle Members: 48.49%"])
```

```
my_circle = plt.Circle( (0,0), 0.5, color= 'white')
```

```
p= plt.gcf()
```


```
p.gca().add_artist(my_circle)
```

```
plt.title("Proportion of Active Members:", fontweight= "bold")
```

```
plt.show()
```

Proportion of Active Members:

Active Members: 51.51%



Data Visualizations 2: Relationships

```
original_df["NumOfProducts"].value_counts(normalize= True) * 100
```

```
1    50.84  
2    45.90  
3     2.66  
4     0.60
```

```
Name: NumOfProducts, dtype: float64
```

```
plt.figure(figsize= (16, 6))
```

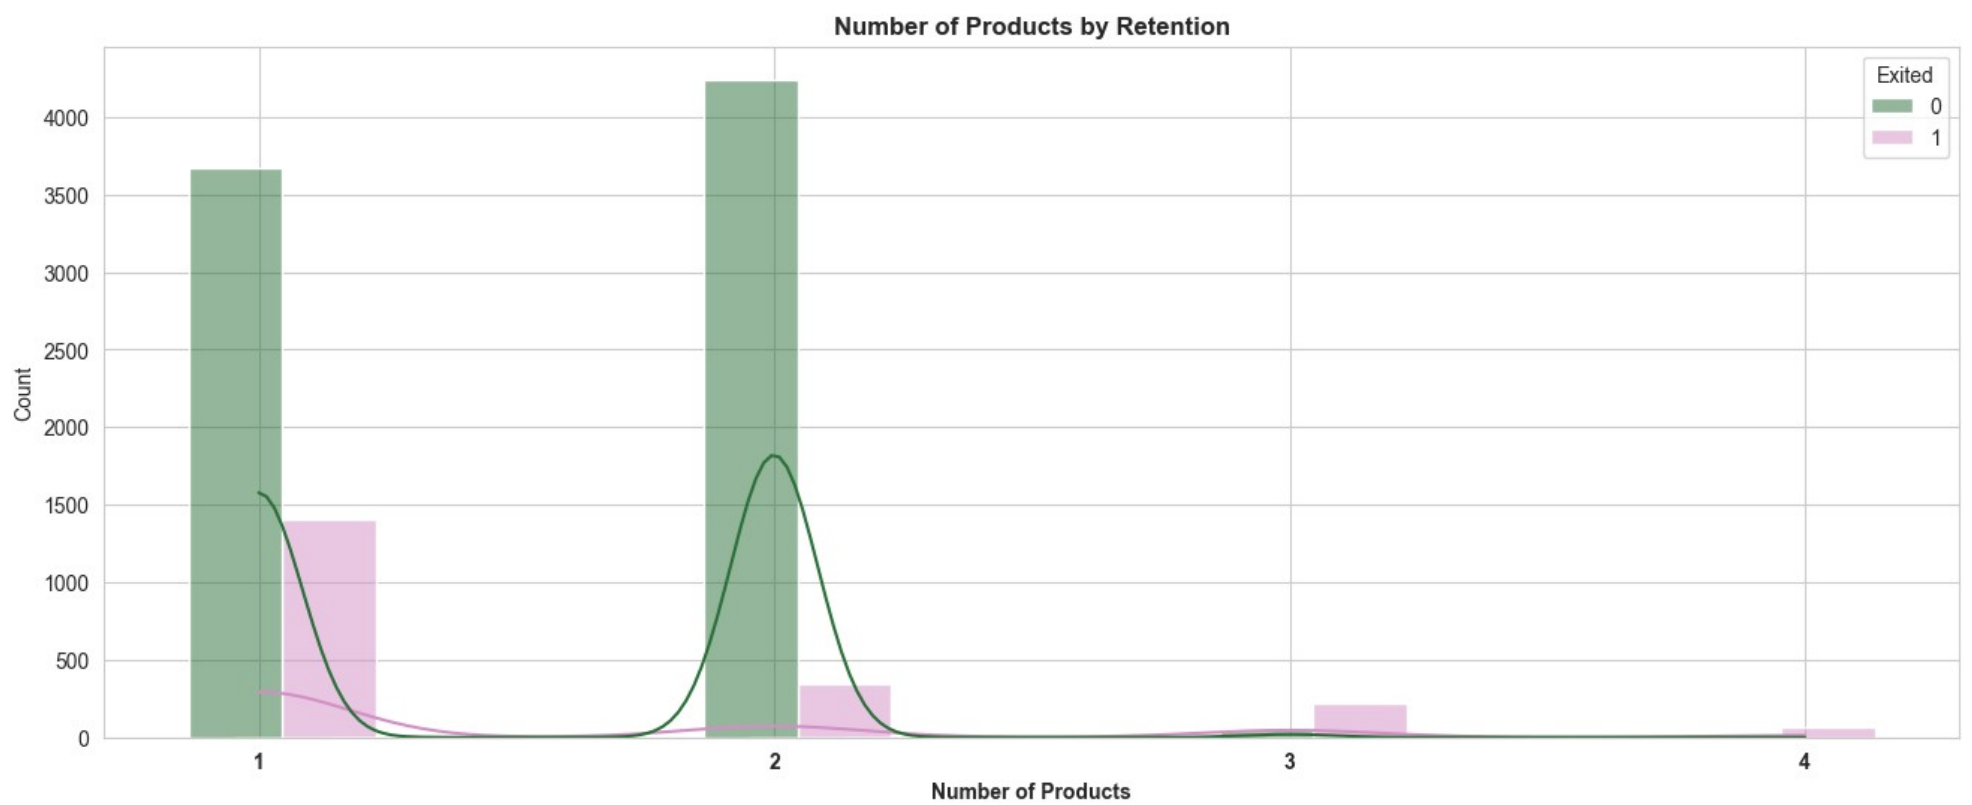
```
sns.histplot(data= original_df, x= "NumOfProducts", shrink= 4, hue= "Exited", multiple= "dodge", palette= "cubehelix", kde= True)
```

```
plt.title("Number of Products by Retention", fontsize= 12, fontweight= "bold")
```

```
plt.xlabel("Number of Products", fontsize= 10, fontweight= "bold")
```

```
plt.xticks([1, 2, 3, 4], fontweight= "bold")
```

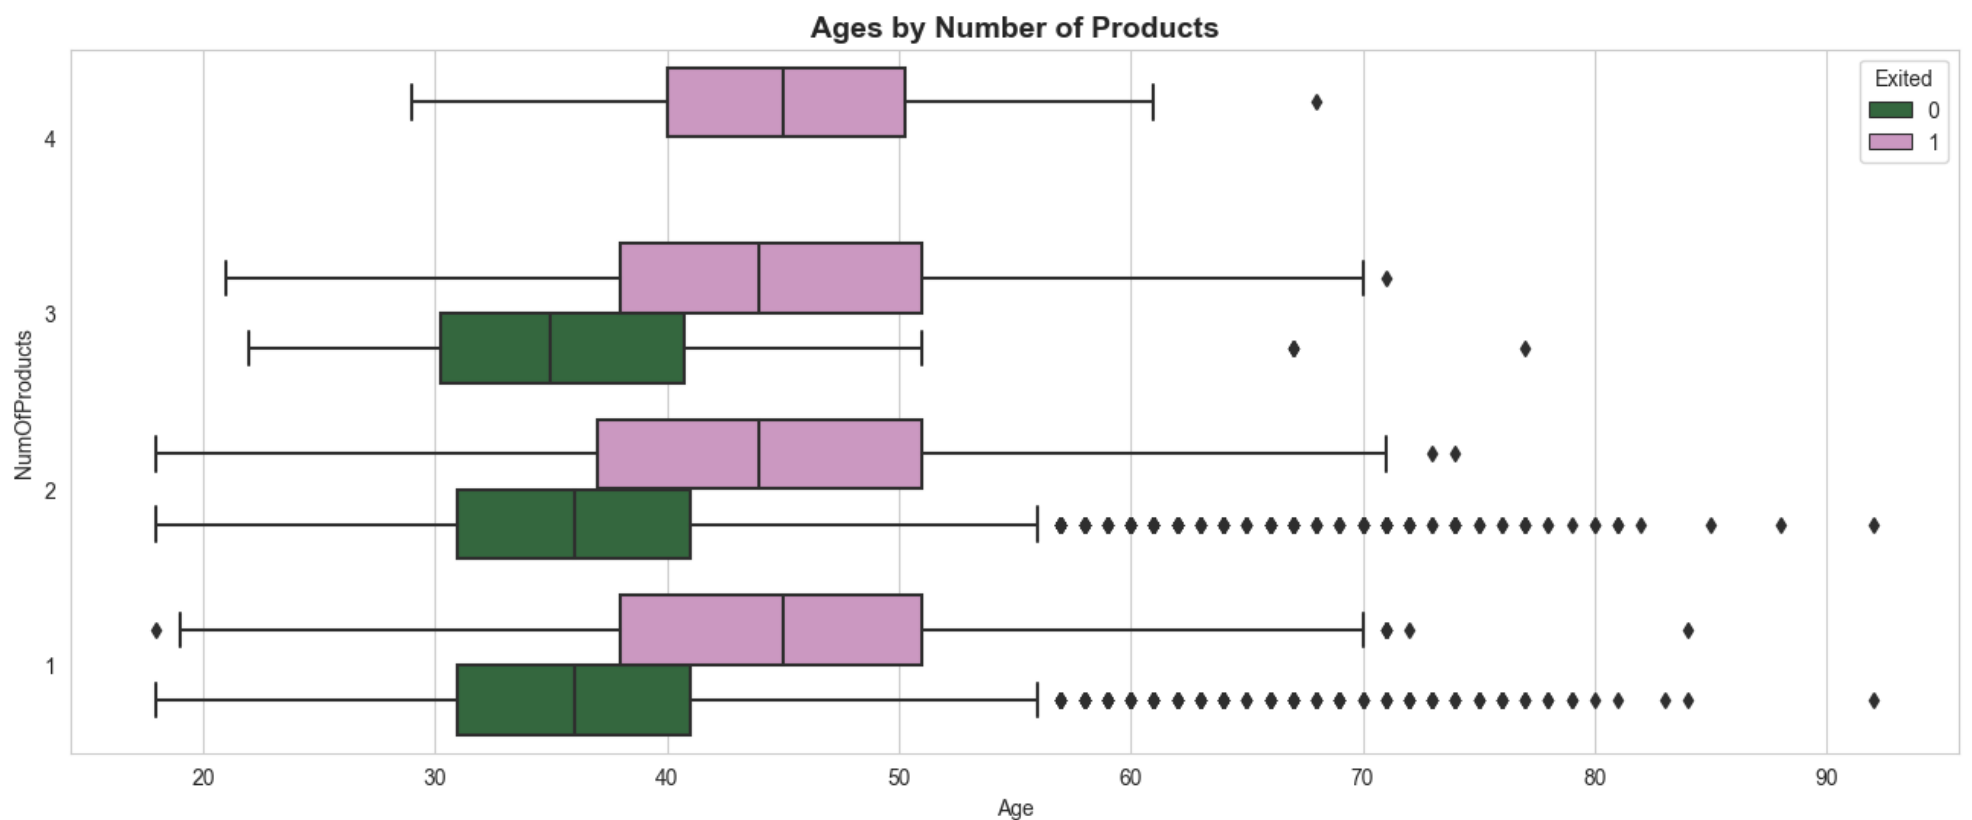
```
plt.show()
```



It's obvious from the above histogram that proportion of customers who have more than 2 products is less than 5% of all customers, and that percentage of churned customers who have 2 products is less than those who have only one product, let's explore than further regarding the customers' age and balance.

```
plt.figure(figsize= (16, 6))
box = sns.boxplot(data= original_df, x= 'Age', y= 'NumOfProducts', hue= 'Exited', orient="h", palette=
"cubehelix")
box.invert_yaxis()

plt.title("Ages by Number of Products", fontsize= 14, fontweight= "bold")
plt.show()
```



From the above boxplots, we conclude that the median age of retained customers is 36 years, while that of churned customers is 45 years.

```
plt.figure(figsize= (16, 6))
box = sns.boxplot(data= original_df, x= 'Balance', y= 'NumOfProducts', hue= 'Exited', orient="h", palette=
"cubehelix")
box.invert_yaxis()

plt.title("Customers' Balance by Number of Products", fontsize= 14, fontweight= "bold")
plt.show()
```

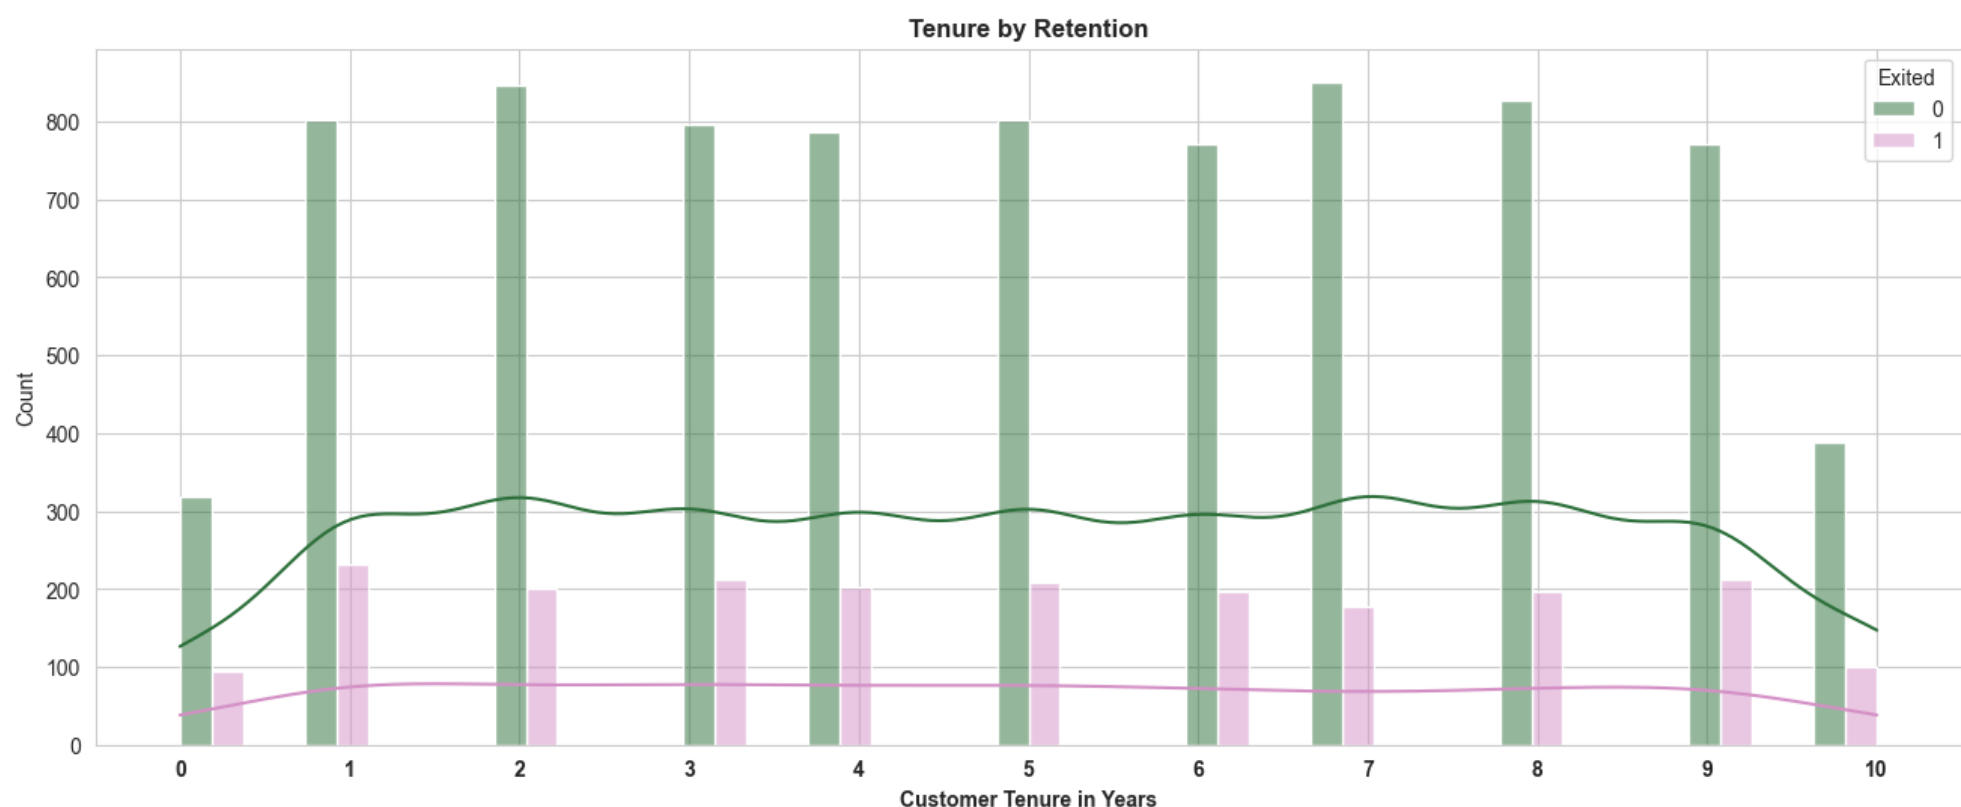



```
original_df["Tenure"].value_counts(normalize= True) * 100
```

```
2    10.48
1    10.35
7    10.28
8    10.25
5    10.12
3    10.09
4     9.89
9     9.84
6     9.67
10    4.90
0     4.13
Name: Tenure, dtype: float64
```

```
plt.figure(figsize= (16, 6))
sns.histplot(data= original_df, x= "Tenure", hue= "Exited", multiple= "dodge", palette= "cubehelix", kde= True)
plt.title("Tenure by Retention", fontsize= 12, fontweight= "bold")

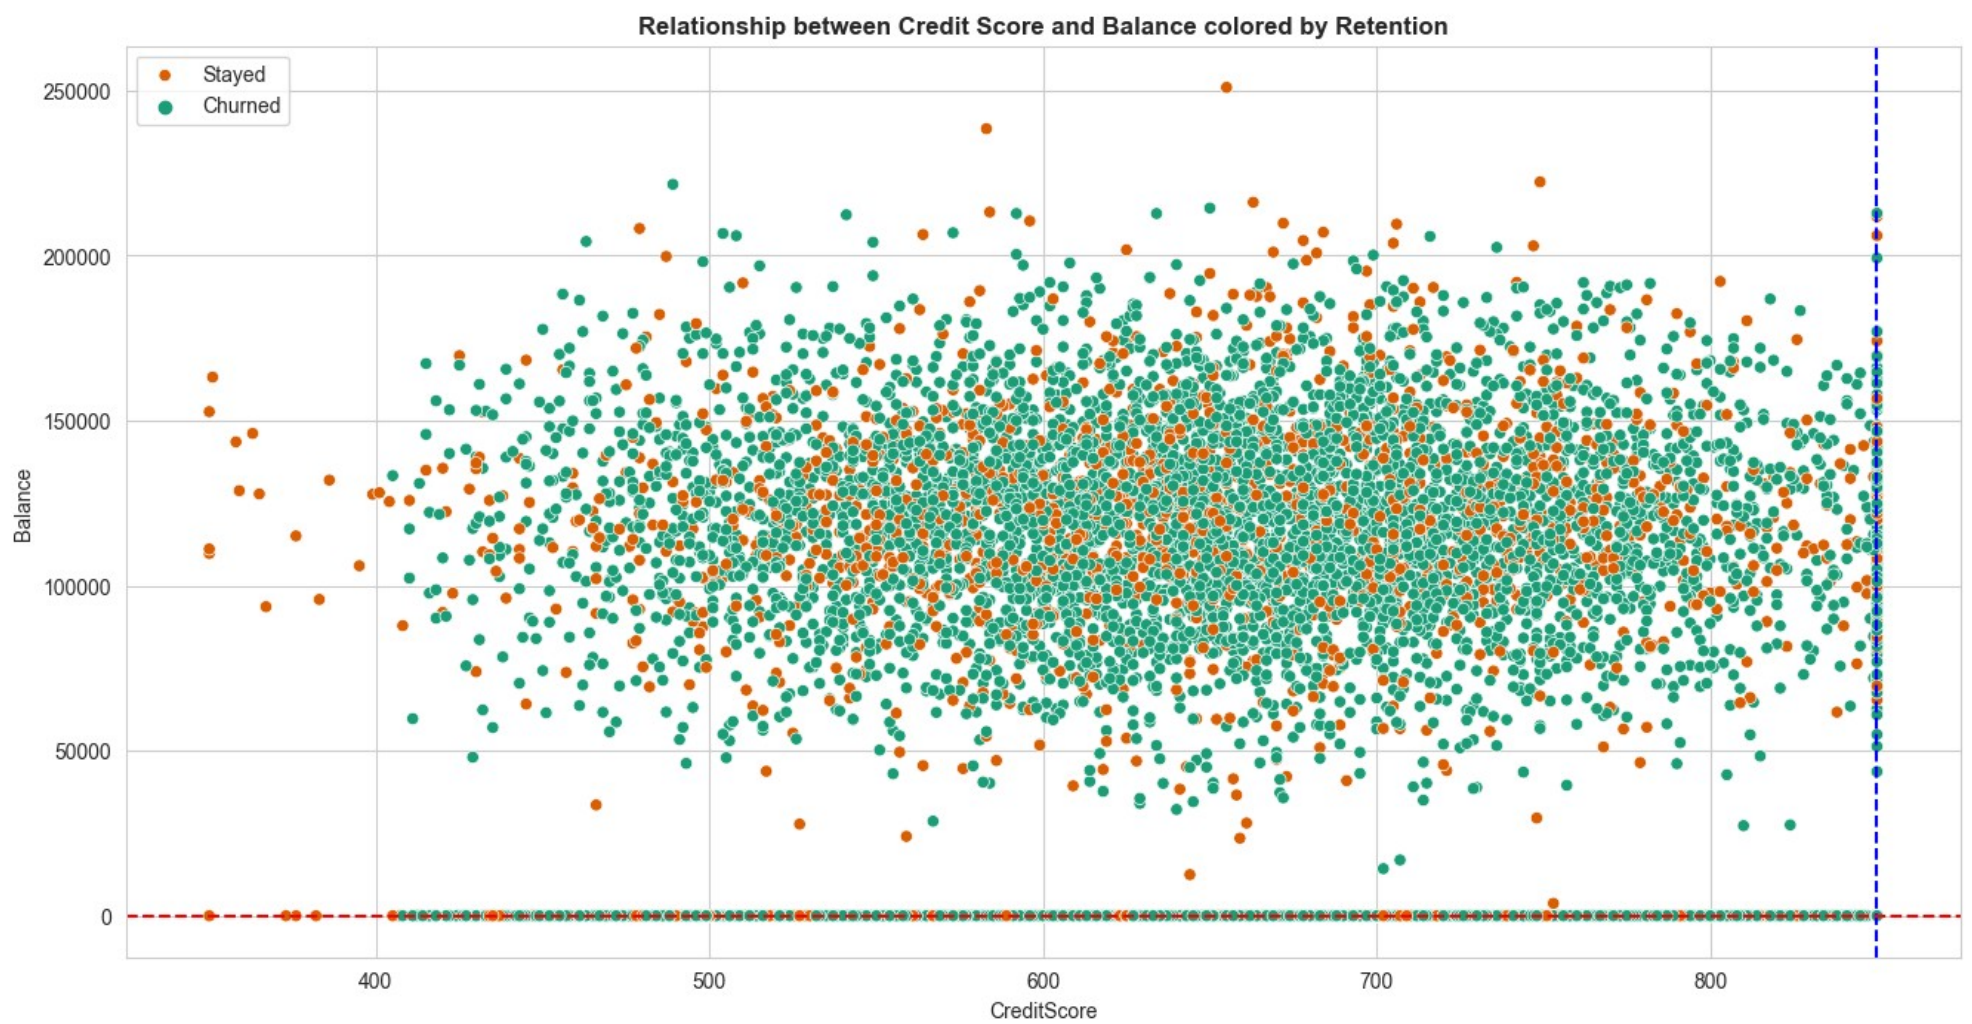
plt.xlabel("Customer Tenure in Years", fontsize= 10, fontweight= "bold")
plt.xticks(range(0, 11), fontweight= "bold")
plt.show()
```



The homogeneous distribution of customers' tenure is obvious, however the proportion of churned customers are roughly stable across all tenures.

```
plt.figure(figsize= (16, 8))
sns.scatterplot(data= original_df, x= "CreditScore", y= "Balance", hue= "Exited")
plt.title("Relationship between Credit Score and Balance colored by Retention", fontsize= 12, fontweight= "bold")

plt.axvline(x= 850, color= "blue", ls= "dashed")
plt.axhline(y= 0, color= "red", ls= "dashed")
plt.legend(["Stayed", "Churned"])
plt.show()
```

From the above scatter plot, we can draw 3 conclusions, which are:

1. Churned customers are distributed across all credit scores and account balance.
2. There are considerable proportion of customers who have a balance of Zero.
3. There are many customers have the full credit score of 850 points.

```
plt.figure(figsize= (16, 8))
sns.scatterplot(data= original_df, x= "EstimatedSalary", y= "Balance", hue= "Exited")
plt.title("Relationship between Estimated Salary and Balance colored by Retention", fontsize= 12, fontweight=
"bold")
```

```
plt.axhline(y= 0, color= "red", ls= "dashed")
plt.legend(["Stayed", "Churned"])
plt.show()
```



The estimated salary variable is distributed very homogeneously across customers, which may be a data collection error or due to data synthesis.

```
plt.figure(figsize= (16, 8))
sns.scatterplot(data= original_df, x= "Age", y= "Balance", hue= "Exited")
plt.title("Relationship between Age and Balance colored by Retention", fontsize= 12, fontweight= "bold")
```

```
plt.axhline(y= 0, color= "red", ls= "dashed")
```



```
plt.legend(["Stayed", "Churned"])
plt.show()
```




We can easily observe from the above scatter plot how proportion of churned customers increases with aging.

```
plt.figure(figsize= (12, 6))

mask_heatmap = np.triu(np.ones_like(original_df.corr(numeric_only= True)))
fig = sns.heatmap(data= original_df.corr(numeric_only= True), annot= True, cmap= "YlGnBu", mask= mask_heatmap)
fig.set_title("Heatmap shows the Pearson Correlation Coefficient among the variables")

plt.xticks(rotation= 45)
plt.show()
```



The heatmap shows a positive correlation between the target variable Exited and each of the Age and the Balance variables, while there is a negative correlation between the target variable and whether the customer is active or not. These assumptions is important before the modelling to understand relationships between the features.

```
#The cost of every churned customer:
```

```
avg_churned_bal = round(original_df[original_df["Exited"] == 1]["Balance"].mean(), 2)
avg_churned_bal
```

91108.54

Feature Engineering:

Feature Extraction:

```
original_df["Loyalty"] = round(original_df["Tenure"] / original_df["Age"] * 100, 2)
original_df["Loyalty"].describe()
```

```
count    10000.000000
mean       13.793457
std        8.950456
min         0.000000
25%        6.450000
50%       12.900000
75%       20.000000
max       55.560000
Name: Loyalty, dtype: float64
```

Feature Selection:

```
churn_df = original_df.drop(["RowNumber", "CustomerId", "Surname", "Gender"], axis= 1)
churn_df.head(5)
```

	CreditScore	Geography	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
0	619	France	42	2	0.00	1	1	
1	608	Spain	41	1	83807.86	1	0	
2	502	France	42	8	159660.80	3	1	
3	699	France	39	1	0.00	2	0	
4	850	Spain	43	2	125510.82	1	1	

	IsActiveMember	EstimatedSalary	Exited	Loyalty
0	1	101348.88	1	4.76
1	1	112542.58	0	2.44
2	0	113931.57	1	19.05
3	0	93826.63	0	2.56
4	1	79084.10	0	4.65

Feature Transformation:

```
churn_df["Geography"].value_counts(normalize= True) * 100
```

```
France    50.14
Germany   25.09
Spain     24.77
Name: Geography, dtype: float64
```

```
churn_df = pd.get_dummies(churn_df)
churn_df.head(5)
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
0	619	42	2	0.00	1	1	
1	608	41	1	83807.86	1	0	
2	502	42	8	159660.80	3	1	
3	699	39	1	0.00	2	0	
4	850	43	2	125510.82	1	1	

	IsActiveMember	EstimatedSalary	Exited	Loyalty	Geography_France	\
0	1	101348.88	1	4.76	1	
1	1	112542.58	0	2.44	0	
2	0	113931.57	1	19.05	1	
3	0	93826.63	0	2.56	1	
4	1	79084.10	0	4.65	0	

	Geography_Germany	Geography_Spain
0	0	0
1	0	1
2	0	0
3	0	0
4	0	1

paCe: Construct Stage

Modelling:

Splitting the data:

```
y = churn_df["Exited"]
```

```
X = churn_df.copy()
X = churn_df.drop("Exited", axis= 1)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size= 0.25, stratify= y, random_state= 42)
```

Scaling the data:

Scaling the data based on their mean and standard deviation, since the naive bayes and logistic regression models requires this scaling to perform properly.

```
#Import the scaler function
from sklearn.preprocessing import StandardScaler
```

```
# Instantiate the scaler
scaler = StandardScaler()
```

```
# Fit the scaler to the training data
scaler.fit(X_train)
```

```
# Scale the training data
X_train = scaler.transform(X_train)
```

```
# Scale the test data
X_test = scaler.transform(X_test)
```

Oversampling the minority class:

Oversampling the class of churned customers from 20% to 60%, since the class imbalance problem often skews the model's predictions to the majority class.

```
# Initial class distribution in the training set
print("Original class distribution in training set:", Counter(y_train))
```

```
# Apply SMOTE to oversample the minority class in the training set
smote = SMOTE(sampling_strategy= 0.6, random_state= 42)
X_resampled, y_resampled = smote.fit_resample(X_train, y_train)
```

```
# New class distribution after resampling
print("Resampled class distribution in training set:", Counter(y_resampled))
```

```
Original class distribution in training set: Counter({0: 5972, 1: 1528})
Resampled class distribution in training set: Counter({0: 5972, 1: 3583})
```

1-Logistic Regression Model Construction:

```
%%time
```

```
lr_scaled = LogisticRegression(max_iter= 500, random_state= 17)
lr_scaled.fit(X_resampled, y_resampled)
```

```
CPU times: total: 0 ns
Wall time: 509 ms
```

```
LogisticRegression(max_iter=500, random_state=17)
```


```
# Get the predictions on test data
y_pred_lr_scaled = lr_scaled.predict(X_test)
```

```
print('F1 Score:', '%.3f' % f1_score(y_test, y_pred_lr_scaled))
print('Precision:', '%.3f' % precision_score(y_test, y_pred_lr_scaled))
print('Recall:', '%.3f' % recall_score(y_test, y_pred_lr_scaled))
print('Accuracy:', '%.3f' % accuracy_score(y_test, y_pred_lr_scaled))
```

```
F1 Score: 0.494
Precision: 0.475
Recall: 0.515
Accuracy: 0.786
```

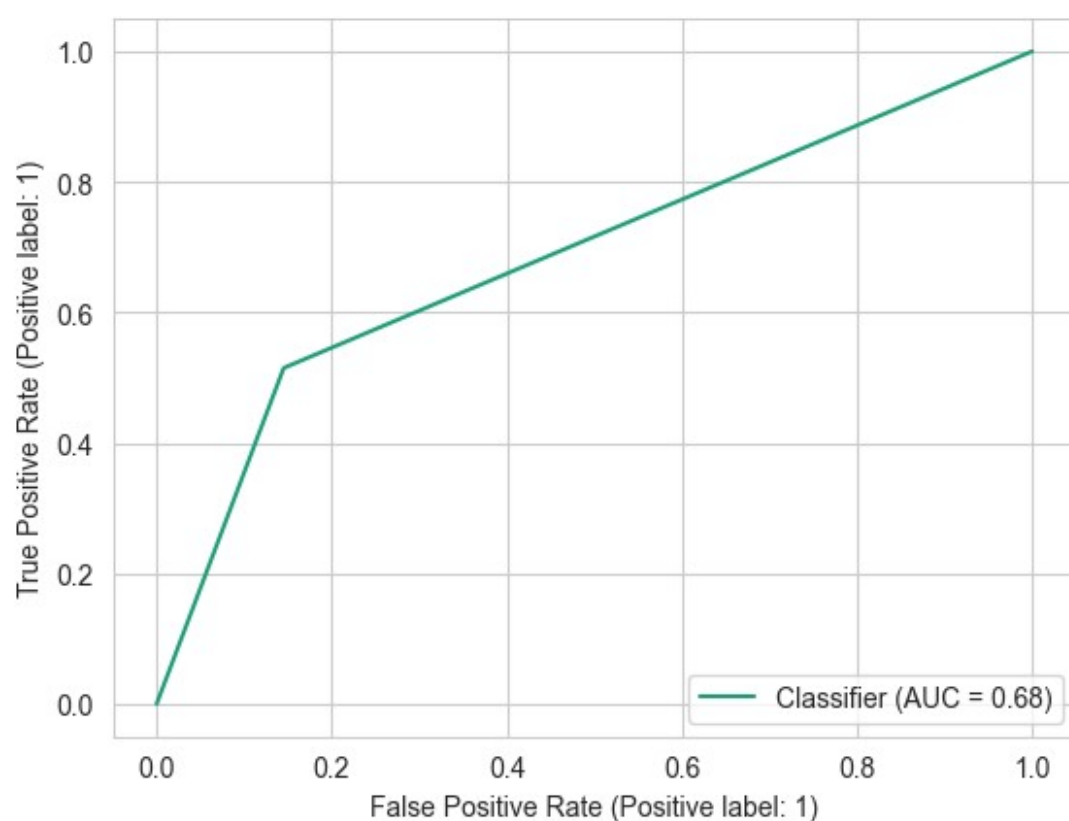
```
def conf_matrix_plot(model, x_data, y_data):
    model_pred = model.predict(x_data)
    cm = confusion_matrix(y_data, model_pred, labels= model.classes_)
    disp = ConfusionMatrixDisplay(confusion_matrix= cm, display_labels= model.classes_)
    disp.plot(values_format= "")
    plt.show()
```

```
conf_matrix_plot(lr_scaled, X_test, y_test)
```



```
RocCurveDisplay.from_predictions(y_test, y_pred_lr_scaled)
plt.show()
```

```
print("roc_auc_score:", round(roc_auc_score(y_test, y_pred_lr_scaled), 3))
```



```
roc_auc_score: 0.685
```

2-Gaussian Naive Bayes Model Construction

```
%%time
```

```
gnb_scaled = GaussianNB()
gnb_scaled.fit(X_resampled, y_resampled)
```

```
CPU times: total: 0 ns
Wall time: 4.45 ms
```

```
GaussianNB()
```


```
# Get the predictions on test data
```

```
y_pred_gnb_scaled = gnb_scaled.predict(X_test)
```

```
print('F1 Score:', '%.3f' % f1_score(y_test, y_pred_gnb_scaled))
print('Precision:', '%.3f' % precision_score(y_test, y_pred_gnb_scaled))
print('Recall:', '%.3f' % recall_score(y_test, y_pred_gnb_scaled))
print('Accuracy:', '%.3f' % accuracy_score(y_test, y_pred_gnb_scaled))
```


```
F1 Score: 0.503
Precision: 0.447
Recall: 0.576
Accuracy: 0.768
```

```
conf_matrix_plot(gnb_scaled, X_test, y_test)
```

```
RocCurveDisplay.from_predictions(y_test, y_pred_gnb_scaled)
plt.show()
```

```
print("roc_auc_score:", round(roc_auc_score(y_test, y_pred_gnb_scaled), 3))
```



```
roc_auc_score: 0.697
```

3-Decision Tree Model Construction

Splitting the data & Oversampling the minority class:

```
y = churn_df["Exited"]
```

```
X = churn_df.copy()
```

```
X = churn_df.drop("Exited", axis= 1)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size= 0.25, stratify= y, random_state= 42)
```

Initial class distribution in the training set

```
print("Original class distribution in training set:", Counter(y_train))
```

Apply SMOTE to oversample the minority class in the training set

```
smote = SMOTE(sampling_strategy= 0.6, random_state= 42)
```

```
X_resampled, y_resampled = smote.fit_resample(X_train, y_train)
```

New class distribution after resampling

```
print("Resampled class distribution in training set:", Counter(y_resampled))
```

```
Original class distribution in training set: Counter({0: 5972, 1: 1528})
```

```
Resampled class distribution in training set: Counter({0: 5972, 1: 3583})
```

```
dt = DecisionTreeClassifier(random_state= 42)
```

```
tree_para = {"max_depth": [3, 4, 5, 6, 7, 8, 9, 10], "min_samples_leaf": [2, 3, 4, 5, 6, 7, 8]}
```

```

scoring = ['accuracy', 'precision', 'recall', 'f1']

dt_cv = GridSearchCV(estimator= dt, param_grid= tree_para, scoring= scoring, cv= 5, refit= 'f1', n_jobs= -1,
verbose= 1)

%%time

dt_cv.fit(X_resampled, y_resampled)

Fitting 5 folds for each of 56 candidates, totalling 280 fits
CPU times: total: 500 ms
Wall time: 5.47 s

GridSearchCV(cv=5, estimator=DecisionTreeClassifier(random_state=42), n_jobs=-1,
             param_grid={'max_depth': [3, 4, 5, 6, 7, 8, 9, 10],
                          'min_samples_leaf': [2, 3, 4, 5, 6, 7, 8]},
             refit='f1', scoring=['accuracy', 'precision', 'recall', 'f1'],
             verbose=1)

print("Best Parameters for the Decision Tree Model:\n", dt_cv.best_params_)

print("\nBest Avgerage Cross-validation F1-score:", "%.3f" % dt_cv.best_score_)

Best Parameters for the Decision Tree Model:
{'max_depth': 9, 'min_samples_leaf': 7}

Best Avgerage Cross-validation F1-score: 0.758

def make_results(model_name, model_object):
    cv_results = pd.DataFrame(model_object.cv_results_)
    best_estimator_results = cv_results.iloc[cv_results["mean_test_f1"].idxmax(), :]
    precision = best_estimator_results.mean_test_precision
    recall = best_estimator_results.mean_test_recall
    f1 = best_estimator_results.mean_test_f1
    accuracy = best_estimator_results.mean_test_accuracy

    table = pd.DataFrame({"Model": [model_name], "F1": [f1], "Precision": [precision], "Recall": [recall],
"Accuracy": [accuracy]})
    return table

dt_cv_results = make_results("Decision_Tree", dt_cv)

dt_cv_results

   Model      F1  Precision  Recall  Accuracy
0  Decision_Tree  0.757611   0.799311   0.725431   0.830141

```

Constructing the optimal Decision Tree Classifier based on the best cross-validation results:

```

decision_tree = DecisionTreeClassifier(max_depth= 9, min_samples_leaf= 7, random_state= 42)
decision_tree.fit(X_resampled, y_resampled)

DecisionTreeClassifier(max_depth=9, min_samples_leaf=7, random_state=42)


y_pred_dt = decision_tree.predict(X_test)

print('F1 Score:', '%.3f' % f1_score(y_test, y_pred_dt))
print('Precision:', '%.3f' % precision_score(y_test, y_pred_dt))
print('Recall:', '%.3f' % recall_score(y_test, y_pred_dt))
print('Accuracy:', '%.3f' % accuracy_score(y_test, y_pred_dt))

F1 Score: 0.597
Precision: 0.596
Recall: 0.597
Accuracy: 0.836

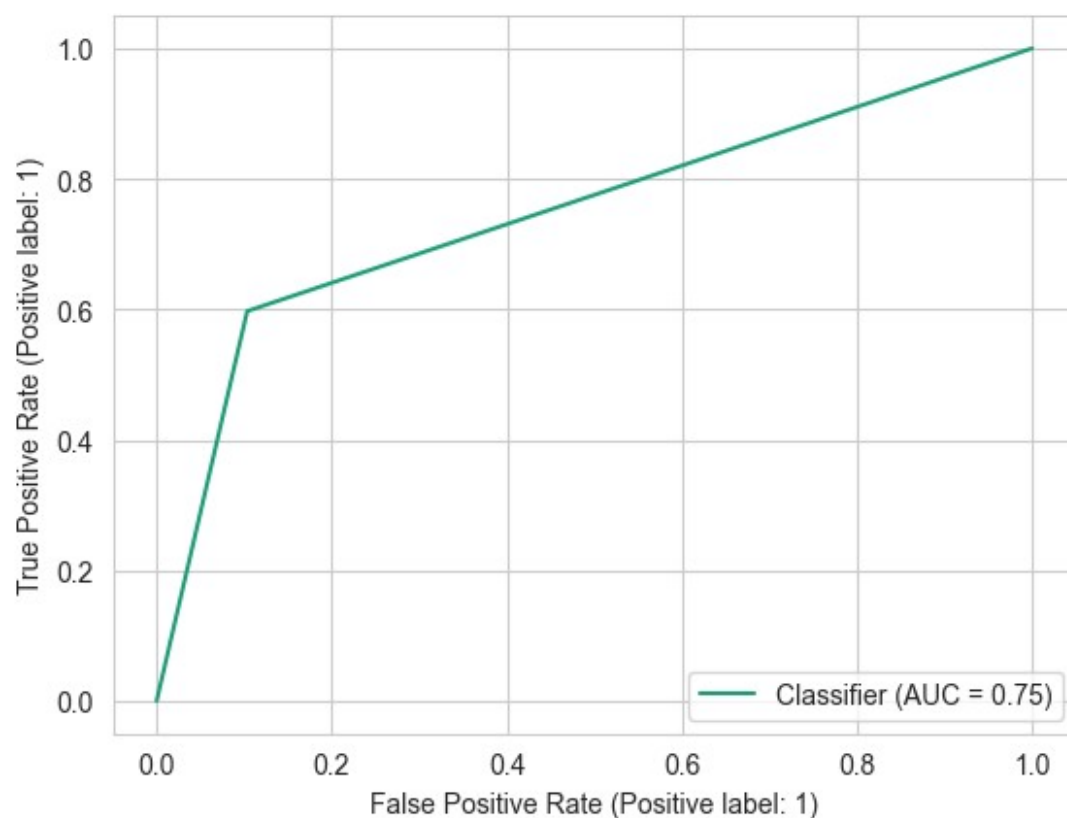
conf_matrix_plot(decision_tree, X_test, y_test)

```



```
RocCurveDisplay.from_predictions(y_test, y_pred_dt)
plt.show()

print("roc_auc_score:", round(roc_auc_score(y_test, y_pred_dt), 3))
```



```
roc_auc_score: 0.747

comparing_models = pd.DataFrame()

comparing_models = comparing_models.append({"Model": "Decision_Tree", "F1": 0.597, "Precision": 0.596,
"Recall": 0.597, "Accuracy": 0.836}, ignore_index= True)

comparing_models = comparing_models.append({"Model": "Naive_Bayes", "F1": 0.503, "Precision": 0.447, "Recall":
0.576, "Accuracy": 0.768}, ignore_index= True)

comparing_models = comparing_models.append({"Model": "Logistic_Regression", "F1": 0.494, "Precision": 0.475,
"Recall": 0.515, "Accuracy": 0.786}, ignore_index= True)
```

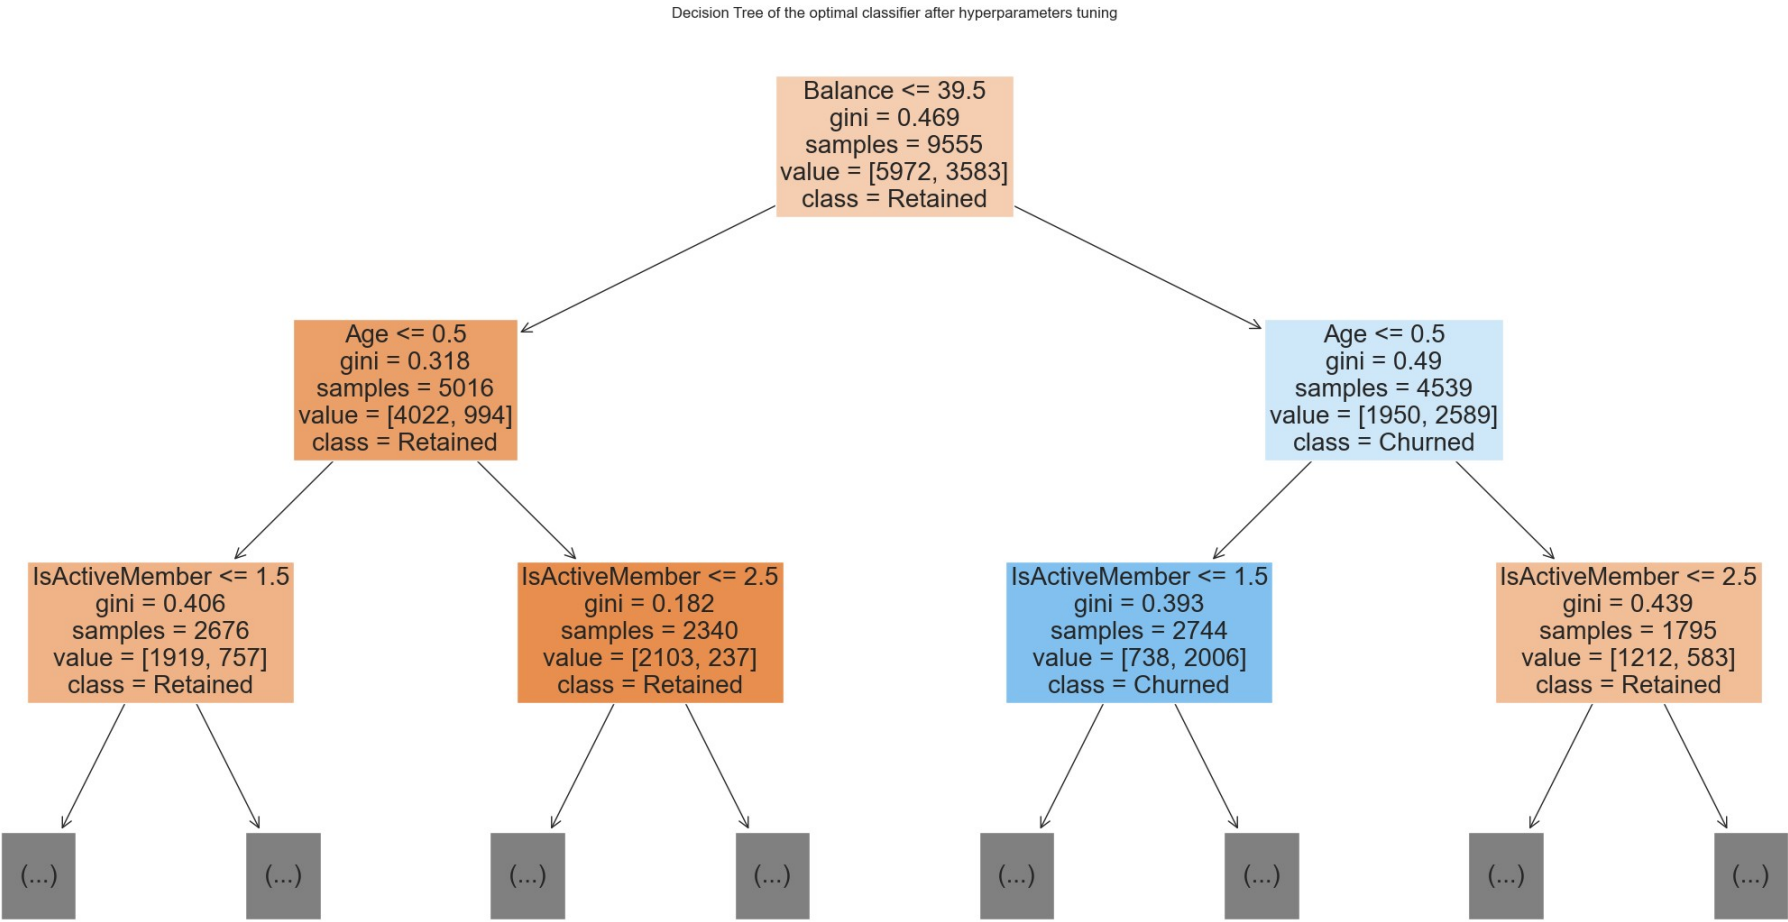
```
comparing_models
```

	Model	F1	Precision	Recall	Accuracy
0	Decision_Tree	0.597	0.596	0.597	0.836
1	Naive_Bayes	0.503	0.447	0.576	0.768
2	Logistic_Regression	0.494	0.475	0.515	0.786

We notice that the F1 Score -which is our benchmark- is higher in the decision tree model (0.525) than the naive bayes model (0.425) and the logistic regression model (0.439), however the False Positives are still higher than the False Negatives despite of the oversampling technique we've conducted earlier.

```
plt.figure(figsize= (28, 14))
plot_tree(decision_tree, max_depth= 2, fontsize= 20,
          feature_names= ['CreditScore', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember',
'EstimatedSalary', 'Age', 'Tenure', 'Loyalty', 'Geography_Germany', 'Geography_Spain', 'Geography_France'],
class_names= ["Retained", "Churned"], filled= True)
```

```
plt.title("Decision Tree of the optimal classifier after hyperparameters tuning")
plt.show()
```



```
importances = decision_tree.feature_importances_ * 100


forest_importances = pd.DataFrame(importances, index= X.columns, columns= ["affect"]).sort_values(by= "affect",
ascending= False)
forest_importances
```

	affect
Age	31.298165
NumOfProducts	17.522073
IsActiveMember	16.185259
Balance	11.096547
Geography_France	6.086130
Geography_Spain	5.919848
Geography_Germany	4.395243
EstimatedSalary	3.054795
Loyalty	2.131416
CreditScore	1.745237
Tenure	0.389989
HasCrCard	0.175299

```
plt.figure( figsize= (10, 5))
plt.pie(x= forest_importances["affect"], labels= ["Age: 31%", "# of Products: 18%", "IsActiveMember: 16%",
"Balance: 11%", "", "", "", "", "", "", "", "", ""])
plt.title("Key Factors affecting Customer Churn by 77%:")
```

```
my_circle = plt.Circle( (0,0), 0.5, color= 'white')
p= plt.gcf()
p.gca().add_artist(my_circle)
plt.show()
```

Key Factors affecting Customer Churn by 77%:



Recommendations for Stakeholders Based on the results of the Decision Tree algorithm, the most effective features on the customer churn rate are:

- 1- Age of the customer.
- 2- Number of Products in which the customer is being involved in.
- 3- Whether the customer is considered as active or not.
- 4- Balance Amount in the customer's account.

The algorithm determined the above features as the most effective ones by approximately 76%, while the other 5 features are not considered effective.

4-Random Forest Model Construction

Cross-validated hyperparameter tuning

```
rf = RandomForestClassifier(random_state= 42)
```

```
cv_params = {'max_depth': [6, 7, 8, 9, 10],
             'min_samples_leaf': [1, 2, 3],
             'min_samples_split': [2, 3, 4],
             'max_features': [2, 3, 4],
             'n_estimators': [75, 100, 125]}
```

```
scoring = ['accuracy', 'precision', 'recall', 'f1']
```

```
rf_cv = GridSearchCV(estimator= rf, param_grid= cv_params, scoring=scoring, cv=5, refit='f1', n_jobs= -1,
verbose= 1)
```

```
#%%time
```

```
#rf_cv.fit(X_resampled, y_resampled)
```

```
path = "D:/Google Advanced Data Analytics/Nuts & Bolts of Machine Learning/Fitted_Models/"
```

```
# Pickle the model
#with open(path+'rf_cv_model.pickle', 'wb') as to_write:
#    pickle.dump(rf_cv, to_write)
```

```
# Read in pickled model
with open(path + 'rf_cv_model.pickle', 'rb') as to_read:
    rf_cv = pickle.load(to_read)
```

```
print("Best Parameters for the Random Forest Model:\n", rf_cv.best_params_)
```

```
print("\nBest Average Cross-validation F1-score:", "%.3f" % rf_cv.best_score_)
```

Best Parameters for the Random Forest Model:

```
{'max_depth': 10, 'max_features': 4, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 125}
```

Best Average Cross-validation F1-score: 0.788

```
rf_cv_results = make_results('Random Forest CV', rf_cv)
rf_cv_results
```

	Model	F1	Precision	Recall	Accuracy
0	Random Forest CV	0.788366	0.852622	0.743864	0.857143

Constructing the optimal Random Forest Classifier based on the best cross-validation results:

```
random_forest = RandomForestClassifier(max_depth= 10, max_features= 4, min_samples_leaf= 1, min_samples_split=
2, n_estimators= 125, random_state= 42)
random_forest.fit(X_resampled, y_resampled)
```




```
RandomForestClassifier(max_depth=10, max_features=4, n_estimators=125,
                        random_state=42)
```

```
y_pred_rf = random_forest.predict(X_test)
```

```
print('F1 Score:', '%.3f' % f1_score(y_test, y_pred_rf))
print('Precision:', '%.3f' % precision_score(y_test, y_pred_rf))
print('Recall:', '%.3f' % recall_score(y_test, y_pred_rf))
print('Accuracy:', '%.3f' % accuracy_score(y_test, y_pred_rf))
```

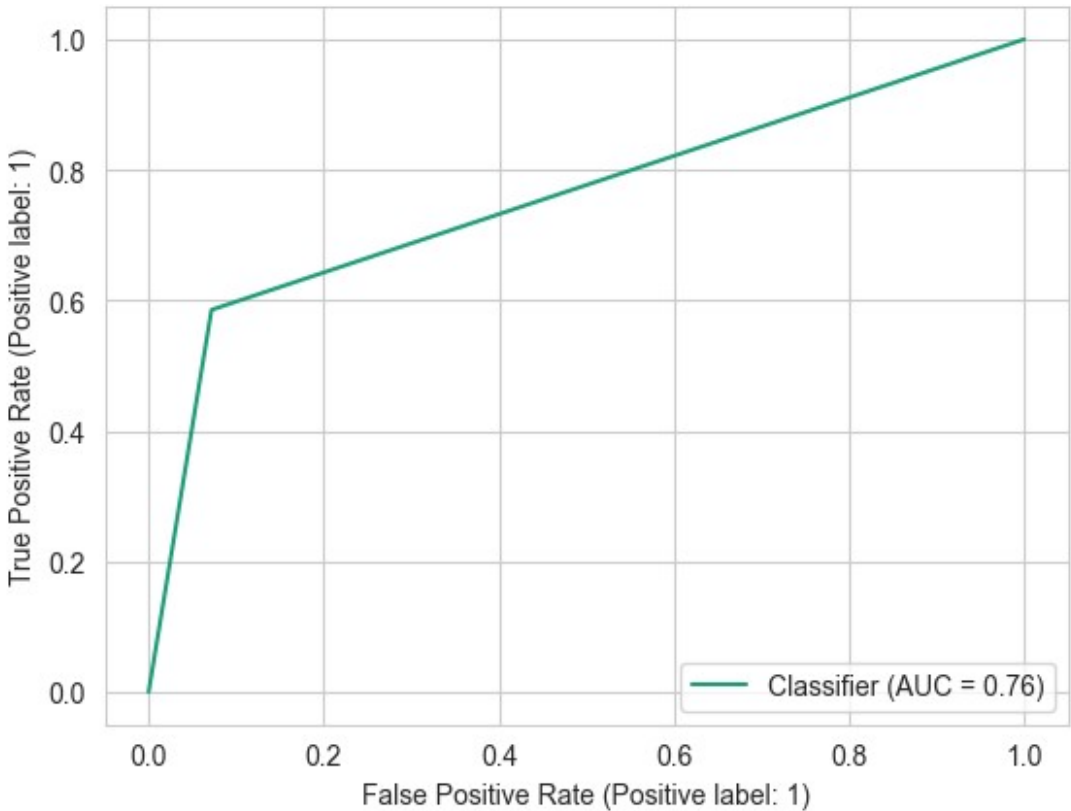
F1 Score: 0.627
Precision: 0.676
Recall: 0.585
Accuracy: 0.858

```
conf_matrix_plot(random_forest, X_test, y_test)
```



```
RocCurveDisplay.from_predictions(y_test, y_pred_rf)
plt.show()
```

```
print("roc_auc_score:", round(roc_auc_score(y_test, y_pred_rf), 3))
```



roc_auc_score: 0.757

```
comparing_models = comparing_models.append({"Model": "Random_Forest", "F1": 0.627, "Precision": 0.676,
"Recall": 0.585, "Accuracy": 0.858}, ignore_index= True)
```

```
comparing_models = comparing_models.sort_values(by= "F1", ascending= False).reset_index(drop= True)
```

```
comparing_models
```

	Model	F1	Precision	Recall	Accuracy
0	Random_Forest	0.627	0.676	0.585	0.858
1	Decision_Tree	0.597	0.596	0.597	0.836
2	Naive_Bayes	0.503	0.447	0.576	0.768
3	Logistic_Regression	0.494	0.475	0.515	0.786


```

feature_importances = random_forest.feature_importances_
feature_importances

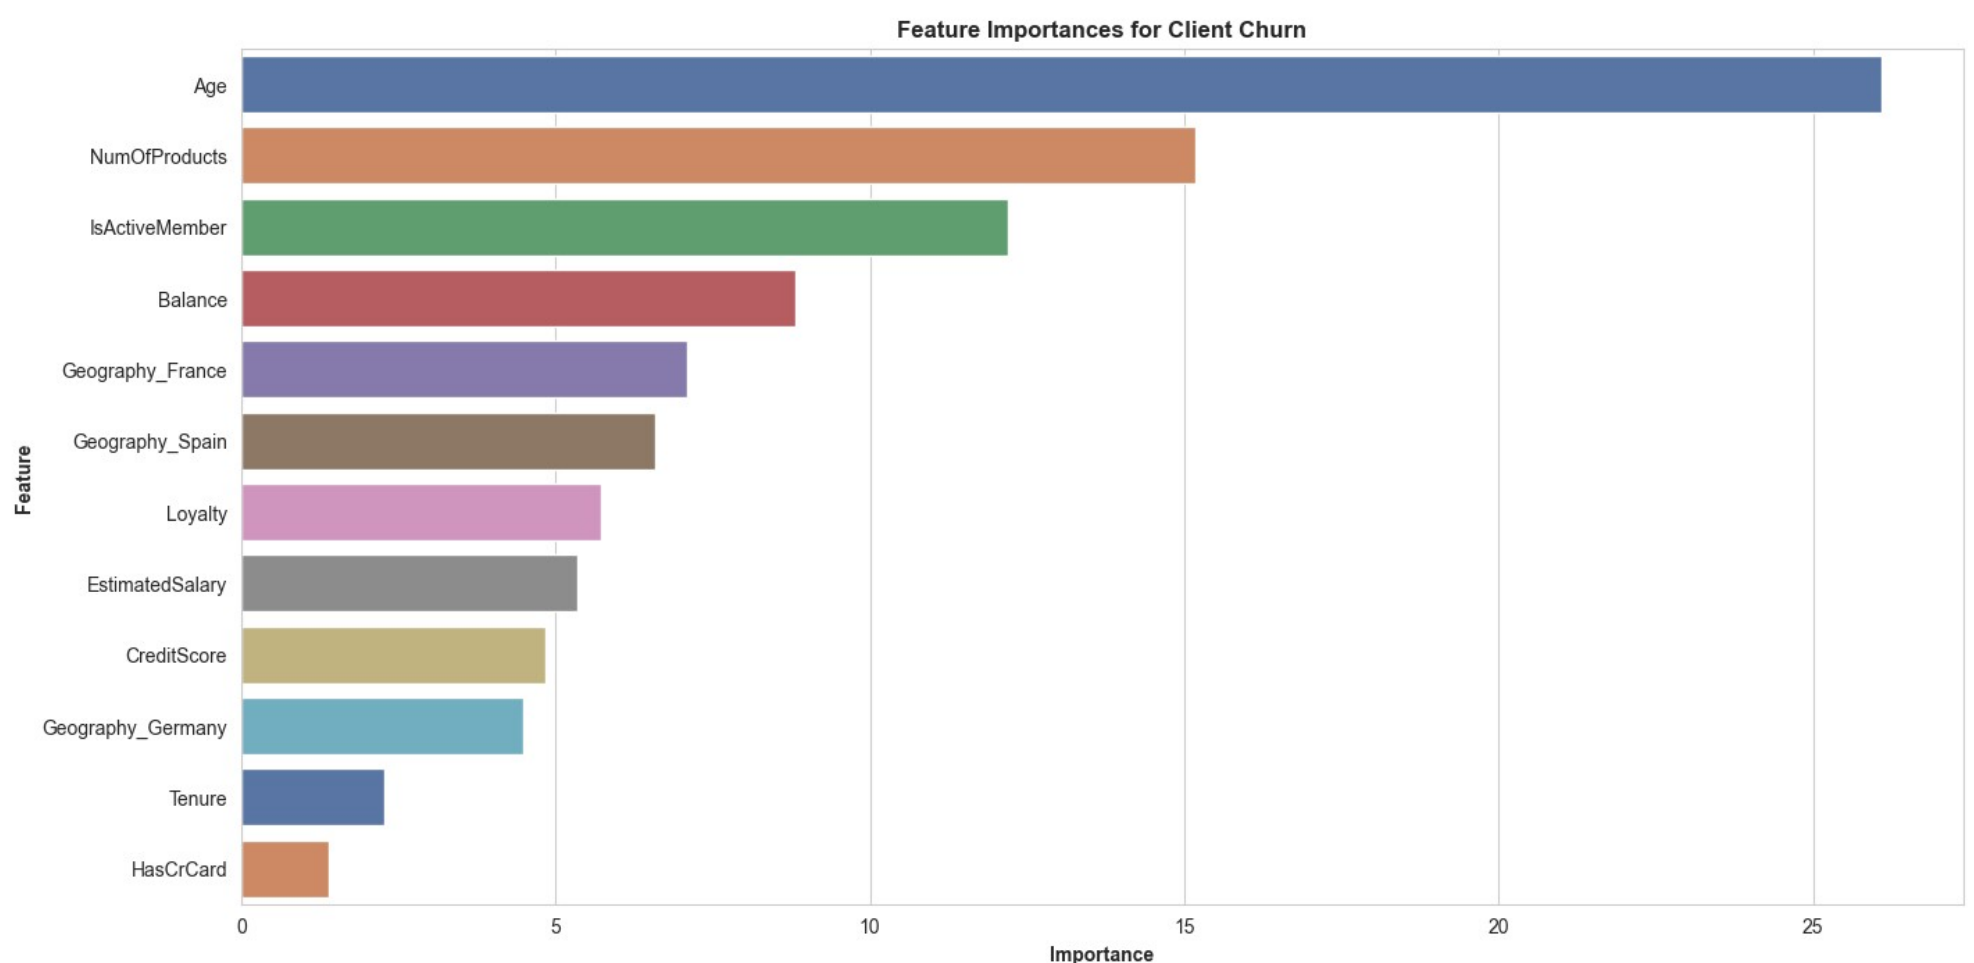
array([0.04835111, 0.26089744, 0.02279709, 0.08810593, 0.15171378,
       0.01385651, 0.12187607, 0.05356716, 0.0571959 , 0.07097223,
       0.04484391, 0.06582285])

forest_importances = pd.DataFrame(list(X.columns), columns= ["feature"])
forest_importances["importance"] = feature_importances * 100
forest_importances = forest_importances[forest_importances["importance"] > 0].sort_values(by="importance",
ascending= False)
forest_importances

   feature  importance
1      Age    26.089744
4  NumOfProducts  15.171378
6  IsActiveMember  12.187607
3      Balance    8.810593
9  Geography_France  7.097223
11 Geography_Spain  6.582285
8      Loyalty    5.719590
7  EstimatedSalary  5.356716
0      CreditScore  4.835111
10 Geography_Germany  4.484391
2      Tenure     2.279709
5     HasCrCard    1.385651

plt.figure(figsize= (16, 8))
sns.barplot(data= forest_importances, x= "importance", y= "feature", palette= "deep", orient= "h")
plt.title("Feature Importances for Client Churn", fontweight= "bold")
plt.ylabel("Feature", fontweight= "bold")
plt.xlabel("Importance", fontweight= "bold")
plt.show()

```



The plot above shows that in this random forest model, age, number of products, member activity and client's balance have the highest importance, in that order. These variables are most helpful in predicting the outcome variable,Exitedt, and they are the same as the onereferred toed by the decision tree model.

5-XGBoost Model Construction

Cross-validated hyperparameter tuning

```

xgb = XGBClassifier(objective= "binary:logistic", random_state=0)

cv_params = {'max_depth': [2, 3, 4, 5, 6], 'min_child_weight': [1, 2, 3, 4, 5], 'learning_rate': [0.05, 0.1, 0.2],
             'n_estimators': [75, 100, 125], 'colsample_bytree': [0.8, 1], 'subsample': [0.8, 1]}

scoring = ['accuracy', 'precision', 'recall', 'f1']

xgb_cv = GridSearchCV(estimator= xgb, param_grid= cv_params, scoring=scoring, cv=5, refit='f1', n_jobs= -1,
verbose= 1)

#%time

#xgb_cv.fit(X_resampled, y_resampled)

```

```

# Pickle the model
#with open(path+'xgb_cv_model.pickle', 'wb') as to_write:
#    pickle.dump(xgb_cv, to_write)

# Read in pickled model
with open(path + 'xgb_cv_model.pickle', 'rb') as to_read:
    xgb_cv = pickle.load(to_read)

print("Best Parameters for the Random Forest Model:\n", xgb_cv.best_params_)

print("\nBest Avgerage Cross-validation F1-score:", "%.3f" % xgb_cv.best_score_)

Best Parameters for the Random Forest Model:
{'colsample_bytree': 0.8, 'learning_rate': 0.2, 'max_depth': 6, 'min_child_weight': 3, 'n_estimators': 100,
'subsample': 0.8}

Best Avgerage Cross-validation F1-score: 0.795

xgb_cv_results = make_results('XGBoost CV', xgb_cv)
xgb_cv_results

```

	Model	F1	Precision	Recall	Accuracy
0	XGBoost CV	0.79467	0.846097	0.763966	0.860492

```

Constructing the optimal XGBoost Classifier based on the best cross-validation results:
xgboost = XGBClassifier(objective = "binary:logistic", max_depth= 6, min_child_weight= 3, learning_rate= 0.2,\
                        n_estimators= 100, colsample_bytree= 0.8, subsample= 0.8, random_state= 42)
xgboost.fit(X_resampled, y_resampled)

XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=0.8, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.2, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=6, max_leaves=None,
              min_child_weight=3, missing=nan, monotone_constraints=None,
              multi_strategy=None, n_estimators=100, n_jobs=None,
              num_parallel_tree=None, random_state=42, ...)

y_pred_xgb = xgboost.predict(X_test)

print('F1 Score:', '%.3f' % f1_score(y_test, y_pred_xgb))
print('Precision:', '%.3f' % precision_score(y_test, y_pred_xgb))
print('Recall:', '%.3f' % recall_score(y_test, y_pred_xgb))
print('Accuracy:', '%.3f' % accuracy_score(y_test, y_pred_xgb))

F1 Score: 0.599
Precision: 0.639
Recall: 0.564
Accuracy: 0.846

conf_matrix_plot(xgboost, X_test, y_test)

```



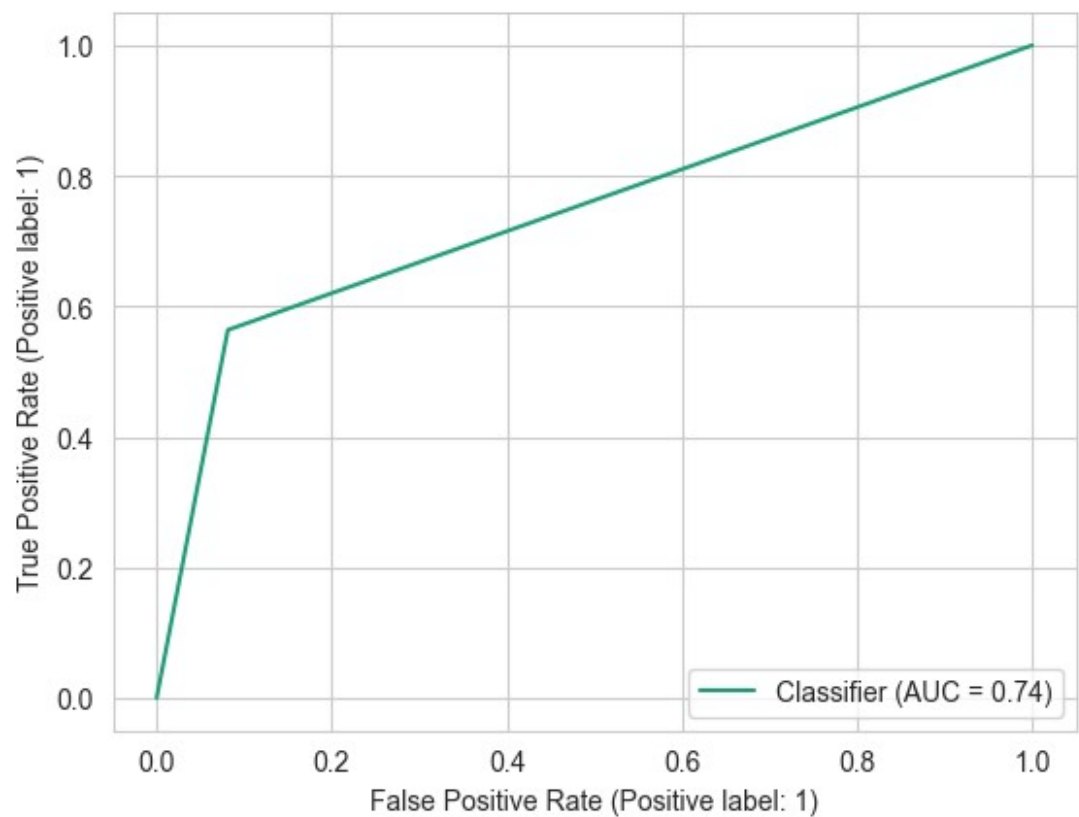
	Predicted 0	Predicted 1	Total
True 0	1829	162	1991
True 1	212	287	500
Total	2041	449	2490

```

RocCurveDisplay.from_predictions(y_test, y_pred_xgb)
plt.show()

print("roc_auc_score:", round(roc_auc_score(y_test, y_pred_xgb), 3))

```



roc_auc_score: 0.741

```
comparing_models = comparing_models.append({"Model": "XGBoost", "F1": 0.599, "Precision": 0.639, "Recall": 0.564, "Accuracy": 0.846}, ignore_index= True)
```

```
comparing_models = comparing_models.sort_values(by= "F1", ascending= False).reset_index(drop= True)
```

comparing_models

	Model	F1	Precision	Recall	Accuracy
0	Random_Forest	0.627	0.676	0.585	0.858
1	XGBoost	0.599	0.639	0.564	0.846
2	Decision_Tree	0.597	0.596	0.597	0.836
3	Naive_Bayes	0.503	0.447	0.576	0.768
4	Logistic_Regression	0.494	0.475	0.515	0.786

pacE: Execute Stage

Comparing the 5 models across the different evaluation metrics based on the testing data:


```
data = {"Model": ["Random_Forest", "Random_Forest", "Random_Forest", "Random_Forest", "XGBoost", "XGBoost", "XGBoost", "XGBoost", "Decision_Tree", "Decision_Tree", "Decision_Tree", "Decision_Tree", "Naive_Bayes", "Naive_Bayes", "Naive_Bayes", "Naive_Bayes", "Logistic_Regression", "Logistic_Regression", "Logistic_Regression", "Logistic_Regression"],\n        "Metric": ["F1", "Precision", "Recall", "Accuracy", "F1", "Precision", "Recall", "Accuracy", "F1", "Precision", "Recall", "Accuracy", "F1", "Precision", "Recall", "Accuracy", "F1", "Precision", "Recall", "Accuracy"],\n        "Score": [0.627, 0.676, 0.585, 0.858, 0.599, 0.639, 0.564, 0.846, 0.597, 0.596, 0.597, 0.836, 0.503, 0.447, 0.576, 0.768, 0.494, 0.475, 0.515, 0.786]}
```

```
comparing_models_ = pd.DataFrame(data)\ncomparing_models_
```

	Model	Metric	Score
0	Random_Forest	F1	0.627
1	Random_Forest	Precision	0.676
2	Random_Forest	Recall	0.585
3	Random_Forest	Accuracy	0.858
4	XGBoost	F1	0.599
5	XGBoost	Precision	0.639
6	XGBoost	Recall	0.564
7	XGBoost	Accuracy	0.846
8	Decision_Tree	F1	0.597
9	Decision_Tree	Precision	0.596
10	Decision_Tree	Recall	0.597
11	Decision_Tree	Accuracy	0.836
12	Naive_Bayes	F1	0.503
13	Naive_Bayes	Precision	0.447
14	Naive_Bayes	Recall	0.576
15	Naive_Bayes	Accuracy	0.768
16	Logistic_Regression	F1	0.494
17	Logistic_Regression	Precision	0.475
18	Logistic_Regression	Recall	0.515
19	Logistic_Regression	Accuracy	0.786

```
plt.figure( figsize= (12, 6))\nfig = sns.barplot(data= comparing_models_, y= "Score", x= "Model", hue= "Metric", palette= "cubehelix", width= 0.7, dodge= True)
```

```
fig.set_title("Comparing the classification models according to the 4 evaluation metrics")
plt.show()
```



We could observe that the Random Forest Model performs the best based on the F1 Score. However, all the models are underrated according to the supposed bank accuracy metrics, so that we can't deploy anyone of them.

Key Takeaways: We need more predictive features to improve the evaluation metrics of our models, since the models perform well on the cross-validation metrics, but they are poorly performing on the testing data.